

Inference-driven resource management and polymorphism in systems programming

Jiří Klepl

MFF UK, September 2022

Introduction

General motivating problem:

Coding low-level systems in simple languages is laborious and error-prone.

Desirable language features:

- **Good generics** to prevent repetition
- **Type inference** to reduce redundant annotations
- **Deep typechecking** to prevent errors
- **Resource management** for runtime correctness

✗ **Unnecessary overhead**

High-level languages and systems programming

General motivating problem:

Coding low-level systems in simple languages is laborious and error-prone.

Desirable language features:

- **Good generics** to prevent repetition
- **Type inference** to reduce redundant annotations
- **Deep typechecking** to prevent errors
- **Resource management** for runtime correctness
- ✗ **Unnecessary overhead**

Language development examples:

- New features of C++
- Rust, Nim, D, Swift, ...
- Ivory (embedded in Haskell)
- C verified by Isabelle/HOL (SeL4 kernel)

Previously in a bachelor thesis: CHM

- “C with type inference”
...based on Hindley-Damas-Milner system
- Overloading and generics
via typeclasses
- Monomorphization
(similar to C++ template instantiation)

Previously in a bachelor thesis: CHM

→ “C with type inference”

...based on Hindley-Damas-Milner system

→ Overloading and generics via typeclasses

→ Monomorphization (similar to C++ template instantiation)

- ✗ Poor inference of
struct member access
(What is the type of `.some_member` ?)
- ✗ No support for resource management
(e.g. locks have to be unlocked, files closed)
- ✗ No subtypes
(C `const`, implicit number conversions)
- ✗ Not much extensibility
(e.g. existentials for run-time polymorphism)
- ✗ Code generation limited to C

Previously in a bachelor thesis: CHM

- “C with type inference”
...based on **Hindley-Damas-Milner system**
- Overloading and generics
via typeclasses
- Monomorphization
(similar to C++ template instantiation)

- ✗ Poor inference of
struct member access
(What is the type of `.some_member` ?)
- ✗ No support for resource management
(e.g. locks have to be unlocked, files closed)
- ✗ No subtypes
(C `const`, implicit number conversions)
- ✗ Not much extensibility
(e.g. existentials for run-time polymorphism)
- ✗ Code generation limited to C

Idea for this thesis: Try a **better type system** and see what happens.

CHM record fields

```
/* CHM required explicit type
   variable declarations like C++
*/
<a : Has_x<a>>
void access(a s) {
    // what type is this:
    s.x;
}
```

CHM (previous work) solution:

- All fields of the same name have to share the same type:

```
struct A { int x; };
struct B { ... x; };
```

The type in B required to be `int` too

- ✗ Field types are still not inferred
- Field access is `struct` → `field`
- We want to overload in both types, and for a given `struct`, we want the `field` type be unique
 - ➔ multi-parameter typeclasses with functional dependencies (MPTCs)

Main results in this thesis

- ✓ Substantial subset of C++ and LLVM assembly output
- ✓ Reduction of code repetition via generics
- ✓ Capable constraint-solver-based type system
 - Support for multiparameter typeclasses allows:
 - ✓ Modeling `struct` member access
 - ✓ Resource management via `Drop` typeclass, which defines the `drop` function
- ✓ Extensible subtype system (C++ constness and data 'kinds')

An example of using generics:

```
// User-provided
//   polymorphic definition
add(auto x, auto y)
{ return (x + y); }

bits32 a1, b1; auto c1;
a1 = 5; b1 = 10;
c1 = add(b1, c1);

bits64 a2, b2; auto c2;
a2 = 5; b2 = 10;
c2 = add(b2, c2);
```

```
// Automatically generated
//   monomorphic definitions
_Madd_$$Lb32...(bits32 x, bits32 y)
{ return (x + y); }

_Madd_$$Lb64...(bits64 x, bits64 y)
{ return (x + y); }

bits32 a1, b1; bits32 c1;
a1 = 5; b1 = 10;
c1 = _Madd_$$Lb32...(a1, b1);

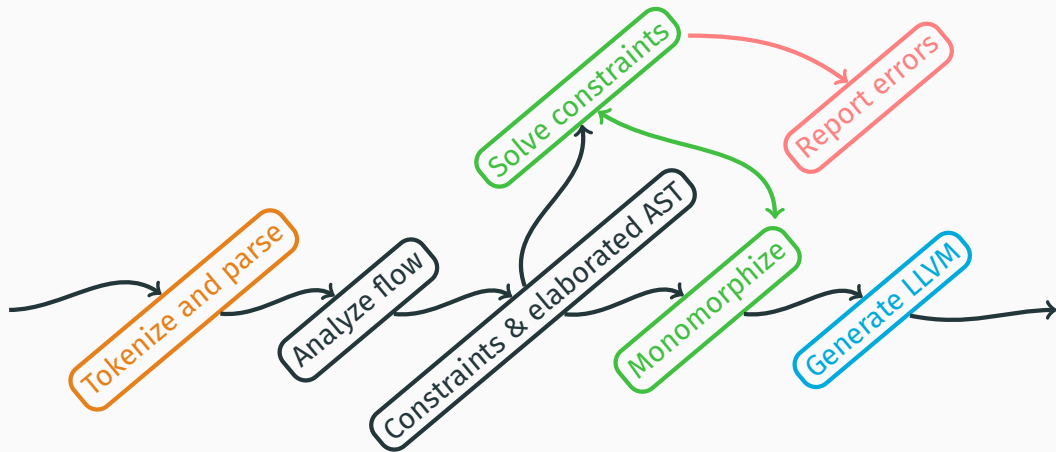
bits64 a2, b2; bits64 c2;
a2 = 5; b2 = 10;
c2 = _Madd_$$Lb64...(a2, b2);
```

Difference between C-- and C

- Non-extended C-- has very primitive type system
 - However, with subtypes that affect register allocation and 'constness' of variables
- C-- contains advanced control-flow primitives
 - Explicit tail-calls: `jump` statements, contrasting `call` statements
 - Interprocedural equivalents to `goto` statements: `cut` `to` statements
- C-- has a simpler structure, for example:
 - No `while` and `for` loops
 - The whole function is one scope, variable declarations hoisted
 - Expressions have no side-effects
- C-- allows for various alignment and aliasing assertions

Compiler

Proof-of-concept compiler



Compiler implemented in Haskell, uses `llvm-hs-pure` and `megaparsec`; **constraint-based type inference** is implemented natively.

- Use the `megaparsec` library
- Generate the AST extending the C++ language syntax specification from:
`https://www.cs.tufts.edu/~nr/c++/extern/man2.pdf`
In-code documentation clearly indicates all deviations from the standard so the implementation could be used for pure C++

Flow analysis – “Flattening” and “Blockifying”

- Splitting into basic-blocks – to facilitate flow analysis and LLVM generation
- Generation of the control-flow graph for each function
- Analysis of variable live-ranges inside the functions
- Insertion of resource management calls to the `drop` function (similar to use of destructors)

Constraint generation and AST elaboration

➡ Input: the AST of the code

➡ Output

- List of constraints to be solved
- The elaborated AST
- The constraints capture the desired type semantics of the code
- Example (simplified):

$x = a + b;$

• $x : \alpha; a : \beta; b : \gamma; a + b : \delta;$

AST elaborations

• $\text{Num } \delta$

the result belongs to the `Num` typeclass (which defines `(+)`)

• $\beta \sim \delta; \gamma \sim \delta$

arguments have the type of the result

• $\alpha \sim \delta$

the assignee and the result share the type

Constraint solving

➡ Input: list of constraints the types have to satisfy

For example:

- Equality constraint $\tau_1 \sim \tau_2$ types τ_1, τ_2 have to be the same
- Class constraint C_τ the type τ has to belong to the typeclass C

➡ Output: assignments of type variables to concrete types, leftover constraints

- We can combine the constraints C_1, C_2 (conjunction)

$$C_1, C_2$$

- And nest the constraint C while binding certain type variables a, b and providing local assumptions A

$$\forall a, b. A \Rightarrow C$$

Monomorphization

- ➡ Input: AST elaborated with the inferred (poly-)types
- ➡ Output: monomorphized AST elaborated with monotypes
 - **Monomorphization** = replacing **polymorphic** definitions with copies with the intended **monotypes**

➡ Input: monomorphized AST

- Elaborated with monotypes
- Split into basic-block structure
- Control-flow and variable liveranges in relation to basic-blocks

➡ Output: the LLVM assembly

- Regular variables translated into LLVM registers
- Each local stack-object `alloca`'ted at function entry
- Translating each basic block separately
 - Variables assigned-to by other basic blocks (or itself) imported via ϕ -nodes
 - Mapping variables to the current values while translating the inner statements

LLVM generation

➡ Input: monomorphized AST

- Elaborated with monotypes
- Split into basic-block structure
- Control-flow and variable liveness in relation to basic-blocks

➡ Output: the LLVM assembly

- Regular variables translated into LLVM registers
- Each local stack-object `alloca`'ed at function entry
- Translating each basic block separately
 - ▼ Variables assigned-to by other basic blocks (or itself) imported via ϕ -nodes
 - Located at the basic block entry
 - For each variable that has to be **alive**
 - Each ϕ -node takes all the possible values of the imported variables
 - The values accompanied with the corresponding exporting basic-block labels
 - Mapping variables to the current values while translating the inner statements

LLVM generation

➡ Input: monomorphized AST

- Elaborated with monotypes
- Split into basic-block structure
- Control-flow and variable liveranges in relation to basic-blocks

➡ Output: the LLVM assembly

- Regular variables translated into LLVM registers
- Each local stack-object `alloca`'ted at function entry
- Translating each basic block separately
 - Variables assigned-to by other basic blocks (or itself) imported via ϕ -nodes
 - ✓ Mapping variables to the current values while translating the inner statements
 - Each statement can change this mapping
 - This mapping then taken as the basic block's export

Details about the type system

- Parametric polymorphism with definable type constructors
- Multi-parameter typeclasses: overloading constrained by a relation
- Automatic calls to resource management functions (bound to a resource lifetime)
- New subtype extension to Hindley-Milner-style type systems

✓ Parametric polymorphism with definable type constructors

- Types: `bits32`, $\alpha \rightarrow \beta$, $\forall a.a$
- Type variables
 - Free (unbound): $\forall a.a \rightarrow \alpha$
can be later *substituted* with a different type
 - Bound: $\forall a.a \rightarrow a$ (id function)
can be instantiated into multiple types according to the given context
- Type constructors
 - Functions: $\tau \rightarrow \sigma$
 - Type constants: `bits32`, `bool`
 - Product types, sum types, and their combination
 - `struct` definitions

Type system details

✓ Multi-parameter typeclasses: overloading constrained by a relation

- Multi-parameter typeclass: a relation over some types (represented by type variables)

```
class Name a b c
```

a typeclass over types a, b, c

```
instance Name bits32 bits32 bits64
```

a typeclass instance given by

```
[a := bits32, b := bits32, c := bits64]
```

- Typeclass method: a function defined under the given typeclass, their type defined via the type variables of the enclosing typeclass

$m :: (a, b) \rightarrow c$: a method with the argument (a, b) returning c

The class instance contains the instance of m typed $(\text{bits32}, \text{bits32}) \rightarrow \text{bits64}$

- Functional dependencies: `class Name a b c | a b -> c` - specifies that, for any instance, each pair a, b uniquely determines c (we can deduce it)

Type system details

- ✓ Automatic calls to resource management functions (bound to a resource lifetime)
 - Closely related to the language's semantics, here we assume C--: a scope is defined by a function activation (from the function entry to one of the corresponding exits); as customary, we tie a resource lifetime to a scope
 - Resource management actions performed by the user-definable `drop` functions
 - usually overloaded in a `Drop` class so that every object issues a different action
 - resource-representing objects defined in `stackdata { ... }` blocks

`Name: someType;` `Name` is a pointer to a stack-allocated object of `someType`

`Name: new someType;` ... and the object represents a resource → `drop`

 - Automatically inserted calls to the `drop` function at function exits (`returns`)
 - Unless such exit follows a `dropped` statement that takes the name of the resource object
 - In C++, for example, scopes with multiple exit points are the minority

✓ New subtype extension to Hindley-Milner-style type systems

- Many systems programming languages have subtypes (`const`, `constexpr`, etc.)
- In the extended C--: *kinds* and *constnesses*
- Kind: a set of registers a value can be stored on
- Constness: specifies when the value is to be computed
compile-time, link-time, run-time
- Each HM-style type is extended with a kind dimension and a constness dimension
- Type (in-)equalities can be specific to a dimension $\tau_1 \leq_{const} \tau_2$
- Different types can have a defined inequality relation in each dimension and yet be incomparable
- Different types can be equal in every dimension and yet still be nonequal

Thank you for attention!

You can try the results on GitHub:

`https://github.com/jiriklepl/masters-thesis-code`

Questions

Backwards-compatibility and extensibility of other languages' type systems

- The extended C++ is not fully backwards-compatible with C++ as it uses stronger type system
 - Compiling C++ programs would be possible with adding some type coercions
 - C++ programs are not meant to be portable between implementations anyway
- The type system as explained in the thesis does not fit dynamic languages like Python, JavaScript or Julia without significantly restricting such languages
 - Extending the type system with existentials as discussed in the thesis would make modelling such languages easier – however, still not complete
- The type system could be modified to fit the needs of a Rust-like language
 - Successful implementation of Rust borrowing through type system would require the type system to be flow-sensitive